

Recurrent Neural Networks (RNN)

Seoul National University

Saerom Park

drsaerompark@gmail.com

References

► Books

- [Goodfellow et al. (2016)] I. Goodfellow, Y. Bengio, and A. Courville, Deep learning, The MIT Press, 2016.

► Online courses

- [Hinton. (2015)] G. Hinton, Online course on Neural Networks for Machine Learning
✓ <https://www.coursera.org/learn/neural-networks>
- [Larochelle. (2014)] H. Larochelle, online course on Neural Networks
✓ http://info.usherbrooke.ca/hlarochelle/neural_networks/content.html
- [Gavves. (2017)] E. Gavves, UVA Deep Learning course
✓ <http://uvadlc.github.io/>
- Borrows slides (some modified) from Larochelle & Gavves

Lecture Overview

- ▶ Recurrent Neural Networks (RNN) for sequences
- ▶ Backpropagation Through Time (bPTT)
- ▶ vanishing and Exploding Gradients and Remedies
- ▶ RNNS using Long Short-Term Memory (LSTM)
- ▶ Applications of RNNS

Sequential Data

► Property

- Next data depend on previous data
- Data inside a sequence are non i.i.d. (identically, independently distributed)

► Example

- The next “word” depends on the previous “words”

► How to deal with sequential data?

- We need **context** and **memory**
- How to model context and memory?

I am Bond , James

Bond


McGuire

Bond

tired

am

!



Sequential Data

► Main task

- Roughly equivalent to predicting what comes next: $\Pr(x) = \prod_i \Pr(x_i | x_1, \dots, x_{i-1})$

► Question


- What about inputs that appear in sequences, such as text?
- Could a neural network handle such modalities?
- How to deal with the different lengths of sequences?

► What is memory?

- Representation of the past
- Information at time step t on a latent space $c(t)$

► Model

- Project information $c(t)$ using parameters θ : $c_{t+1} = h(x_{t+1}, c_t; \theta)$
- Share parameters θ for all time steps: $c_{t+1} = h(x_{t+1}, h(x_t, h(x_{t-1}, \dots, h(x_1, c_0; \theta); \theta); \theta); \theta)$




Memory

► What is memory?

- Representation of the past
- Information at time step t on a latent space $c(t)$

► Model

- Project information $c(t)$ using parameters θ : $c_{t+1} = h(x_{t+1}, c_t; \theta)$
- Share parameters θ for all time steps: $c_{t+1} = h(x_{t+1}, h(x_t, h(x_{t-1}, \dots, h(x_1, c_0; \theta); \theta); \theta); \theta)$



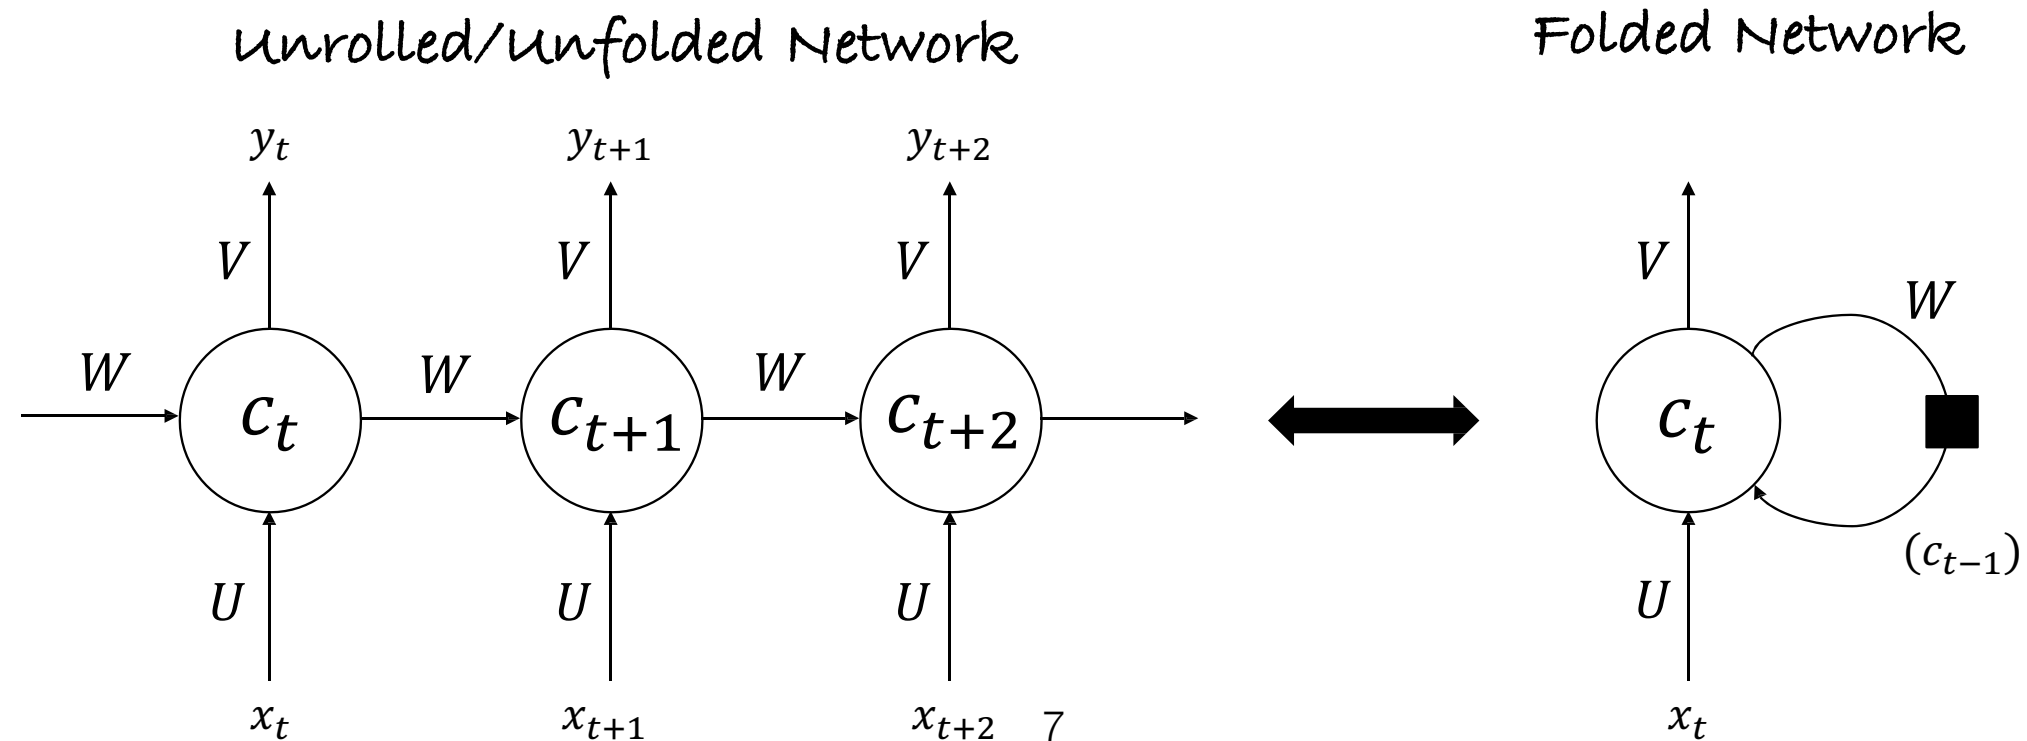
Recurrent Neural Network

► Folding the memory

- Simplify the repeated parameters and structure

► only two equations

- Memory (cell) $c_t = \tanh(b + U x_t + W c_{t-1})$
- output $y_t = \text{softmax}(a + V c_t)$



Introducing Sequential Data

► The different categories of sequence modeling

- Many-to-one: 입력 데이터는 sequence이지만 출력은 고정 크기의 벡터
✓ Sentiment Analysis
- One-to-many: 입력 데이터는 고정 크기의 벡터, 출력은 sequence
✓ Image captioning
- Many-to-many: 입력 및 출력이 모두 sequence
✓ eg. video classification, translation



RNN Example

► component sizes

- Input size: 5
- memory size: 3 (3 units)
- Output size: 10

► Parameters

- cell size ($c(t)$)?
- Input projection (U)?
- cell projection (W)?
- Output projection (V)?

$$c_t = \tanh(U x_t + W c_{t-1})$$

$$y_t = \text{softmax}(V c_t)$$

RNN VS. MLP?

▶ what is really different?

- 입력 (input) 데이터가 시계열 또는 일반적인 시퀀스 데이터임
- 현재 시점의 예측 값을 계산할 때에 이전 시간의 정보를 함께 사용함
- RNN은 입력 데이터 자체가 시퀀스 데이터이고 마지막 시점에서만 output을 예측하도록 구성하는 것도 가능함

▶ RNN 학습

- MLP의 Backpropagation과 본질적으로 동일
- 차이는 RNN의 구조가 시간에 따라 연결되어 있기 때문에 Backprop이 시간을 거슬러 올라가며 적용된다는 것 뿐임

Training RNNS

► Loss Function

- classification: cross-entropy loss

$$P = \prod_{t,k} y_{tk}^{l_{tk}} \Rightarrow \mathcal{L} = -\log P = \sum_t \mathcal{L}_t = -\frac{1}{T} \sum_t \sum_k l_{tk} \log y_{tk}$$

► Backpropagation Through Time (BPTT)

- Again, chain rule!
- only difference \mathcal{L} gradients survive over time steps
- we have three parameter (w, v, u) : $\frac{\partial \mathcal{L}}{\partial V}, \frac{\partial \mathcal{L}}{\partial W}, \frac{\partial \mathcal{L}}{\partial U}$

Backpropagation Through Time: An Example

► Gradients for parameters


- Step by step explanation

✓ <http://www.wildml.com/2015/10/recurrent-neural-networks-tutorial-part-3-backpropagation-through-time-and-vanishing-gradients/>

$$\begin{aligned} c_t &= \tanh(b + Ux_t + Wc_{t-1}) \\ y_t &= \text{softmax}(a + Vc_t) \\ \mathcal{L} &= -\sum_t l_t \log y_t = \sum_t \mathcal{L}_t \end{aligned}$$

$$\frac{\partial \mathcal{L}}{\partial c_t} = W^T \text{diag}(1 - c_{t+1}^2) \frac{\partial \mathcal{L}}{\partial c_{t+1}} + V^T (y_t - l_t) \quad \frac{\partial \mathcal{L}_\tau}{\partial c_\tau} = V^T (y_\tau - l_\tau)$$

$$P = \prod_{t,k} y_{tk}^{l_{tk}} \Rightarrow \mathcal{L} = -\log P = \sum_t \mathcal{L}_t = -\frac{1}{T} \sum_t \sum_k l_{tk} \log y_{tk}$$



$$\frac{\partial \mathcal{L}}{\partial U} = \sum_{t=1}^{\tau} \frac{\partial \mathcal{L}}{\partial U_t} = \sum_{t=1}^{\tau} \frac{\partial \mathcal{L}}{\partial c_t} \left(\frac{\partial c_t}{\partial U_t} \right)^T = \sum_{t=1}^{\tau} \frac{\partial \mathcal{L}}{\partial c_t} \cdot \text{diag}(1 - c_t^2) \cdot x_t^T$$

$$\frac{\partial \mathcal{L}}{\partial W} = \sum_{t=1}^{\tau} \frac{\partial \mathcal{L}}{\partial W_t} = \sum_{t=1}^{\tau} \frac{\partial \mathcal{L}}{\partial c_t} \left(\frac{\partial c_t}{\partial W_t} \right)^T = \sum_{t=1}^{\tau} \frac{\partial \mathcal{L}}{\partial c_t} \cdot \text{diag}(1 - c_t^2) \cdot c_{t-1}^T$$

$$\frac{\partial \mathcal{L}}{\partial b} = \sum_{t=1}^{\tau} \frac{\partial \mathcal{L}}{\partial b_t} = \sum_{t=1}^{\tau} \left(\frac{\partial c_t}{\partial b_t} \right)^T \frac{\partial \mathcal{L}}{\partial c_t} = \sum_{t=1}^{\tau} \text{diag}(1 - c_t^2) \cdot \frac{\partial \mathcal{L}}{\partial c_t}$$

$$\frac{\partial \mathcal{L}}{\partial V} = \sum_{t=1}^{\tau} \frac{\partial \mathcal{L}}{\partial \mathcal{L}_t} \frac{\partial \mathcal{L}_t}{\partial \alpha_t} \frac{\partial \alpha_t}{\partial V} = (l_t - y_t) \cdot \frac{\partial (Vc_t)}{\partial V} = \sum_{t=1}^{\tau} (y_t - l_t) \cdot c_t^T$$

$$\frac{\partial \mathcal{L}}{\partial a} = \sum_{t=1}^{\tau} \frac{\partial \mathcal{L}}{\partial \mathcal{L}_t} \left(\frac{\partial \alpha_t}{\partial a} \right)^T \frac{\partial \mathcal{L}_t}{\partial \alpha_t} = \sum_{t=1}^{\tau} I \cdot (y_t - l_t) = \sum_{t=1}^{\tau} (y_t - l_t)$$

challenge of Long-term Dependencies

► vanishing gradients

- After a few time steps the gradients become almost 0

► Exploding gradients

- After a few time steps the gradients become huge

► can't capture long-term dependencies

- To make it simpler, assume that $c_t = W \cdot c_{t-1}$ and W admits an Eigen-decomposition of the form $W = V\Lambda V^{-1}$, then


$$c_t = W \cdot c_{t-1} = W^t \cdot c_0 = V\Lambda^t V^{-1} \cdot c_0$$

- ✓ The eigenvalues are raised to the power of t causing eigenvalues with magnitude less than one to decay to zero and eigenvalues with magnitude greater than one to explode. Any component of c_0 that is not aligned with the largest eigenvector will eventually be discarded.

How to solve these problem?

► Exploding gradients

- Scale the gradients to a threshold




► vanishing gradients

- It can make long-term dependencies negligible
- Learning focuses on the short-term only
- Difficult to detect!

$$\frac{\partial \mathcal{L}_t}{\partial W} = \sum_{s=1}^t \frac{\partial \mathcal{L}_r}{\partial c_t} \frac{\partial y_t}{\partial c_t} \frac{\partial c_t}{\partial c_s} \frac{\partial c_s}{\partial W}$$
$$\frac{\partial c_t}{\partial c_s} = \prod_{t \geq k \geq s} \frac{\partial c_k}{\partial c_{k-1}} = \prod_{t \geq k \geq s} W \cdot \partial \tanh(c_{k-1})$$

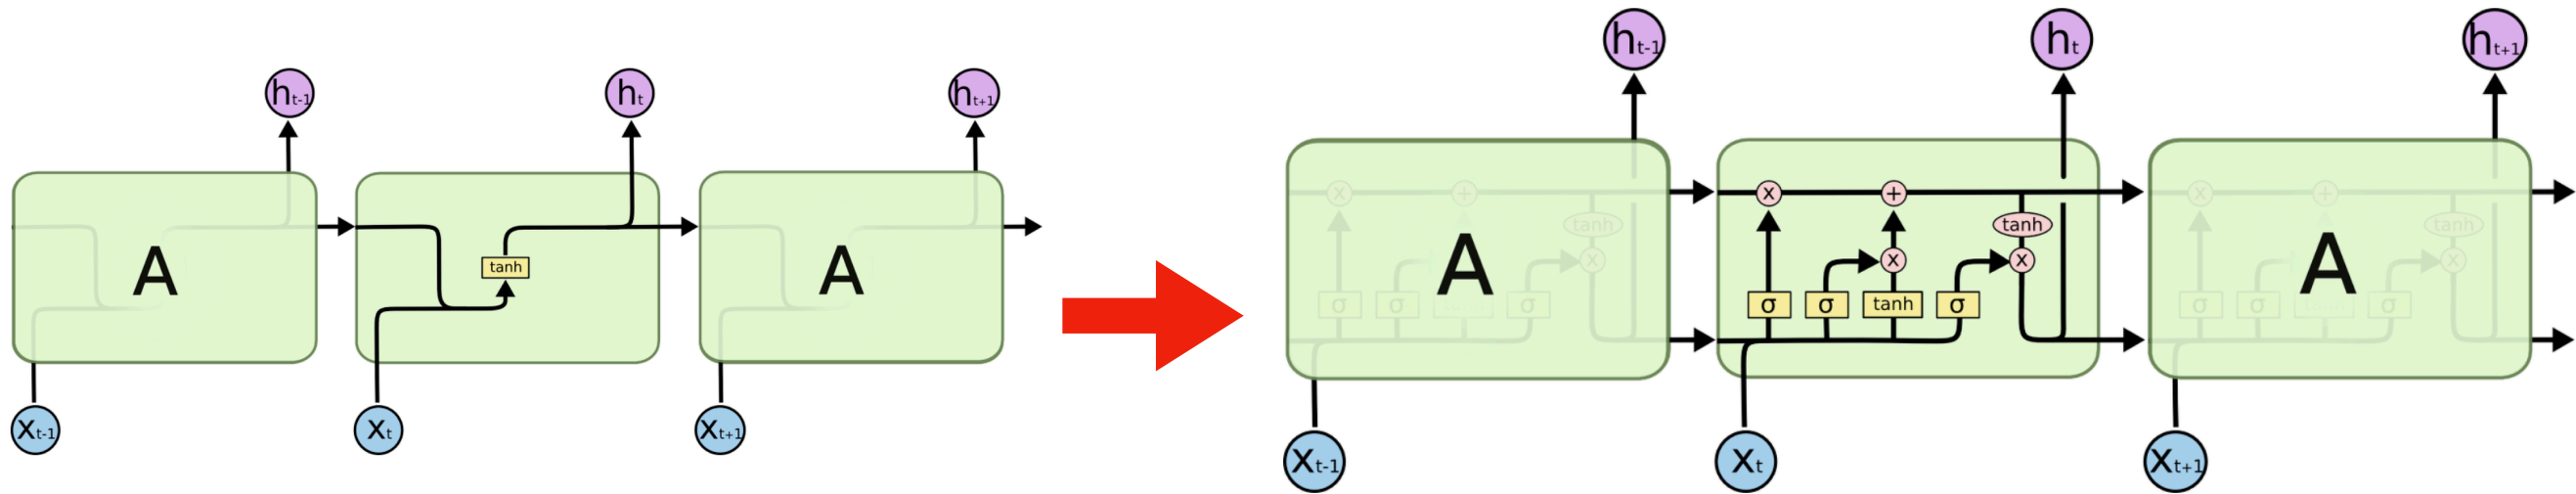
► Advanced RNNS

- Gradient 연산들이 곱셈이 아닌 더하기 연산으로 구성된 RNNS
- Long-short term memory (LSTM) module
- Gated recurrent unit (GRU) module




LSTM network

▶ RNN vs. LSTM structure



- LSTM 은 오차의 gradient 가 시간을 거슬러서 잘 흘러갈 수 있도록 도와줌
- Backprop 과정에서 오차의 값이 더 잘 유지되도록 함으로 long-term 정보도 더 잘 기억하도록 함
- 일반적인 RNN units 은 곱하기로만 이루어져 있는 반면 피드백을 더하기로 이용으로써 sigmoid 곱에 대한 gradient vanishing 문제 해결

▶ 기호



LSTM components

▶ 저장할 정보 (cell states)

$$C_t = f_t * C_{t-1} + i_t * \tilde{C}_t$$

▶ 새로운 정보

- Forget gate layer: cell state에서 어떤 정보를 버릴지 (0~1)

$$f_t = \sigma(W_f \cdot [h_{t-1}, x_t] + b_f)$$

- Input gate layer: 어떤 값을 업데이트 할지 결정 (0~1)

$$i_t = \sigma(W_i \cdot [h_{t-1}, x_t] + b_i)$$

- cell state 에 더해질 수 있는 새로운 후보값 (-1~1)

$$\tilde{C}_t = \tanh(W_C \cdot [h_{t-1}, x_t] + b_C)$$

▶ 출력

- Output gate layer: 어떤 출력값을 출력할지 결정 (0~1)

$$o_t = \sigma(W_o \cdot [h_{t-1}, x_t] + b_o)$$

- 실제 출력값

$$h_t = o_t * \tanh(C_t)$$

LSTM components

▶ 저장할 정보 (cell states)

$$C_t = f_t * C_{t-1} + i_t * \tilde{C}_t$$

▶ 새로운 정보

- Forget gate layer: cell state에서 어떤 정보를 버릴지 (0~1)

$$f_t = \sigma(W_f \cdot [h_{t-1}, x_t] + b_f)$$

- Input gate layer: 어떤 값을 업데이트 할지 결정 (0~1)

$$i_t = \sigma(W_i \cdot [h_{t-1}, x_t] + b_i)$$

- cell state 에 더해질 수 있는 새로운 후보값 (-1~1)

$$\tilde{C}_t = \tanh(W_C \cdot [h_{t-1}, x_t] + b_C)$$


▶ 출력

- Output gate layer: 어떤 출력값을 출력할지 결정 (0~1)

$$o_t = \sigma(W_o [h_{t-1}, x_t] + b_o)$$

- 실제 출력값

$$h_t = o_t * \tanh(C_t)$$



LSTM components

▶ 저장할 정보 (cell states)

$$C_t = f_t * C_{t-1} + i_t * \tilde{C}_t$$

▶ 새로운 정보

- Forget gate layer: cell state에서 어떤 정보를 버릴지 (0~1)

$$f_t = \sigma(W_f \cdot [h_{t-1}, x_t] + b_f)$$

- Input gate layer: 어떤 값을 업데이트 할지 결정 (0~1)

$$i_t = \sigma(W_i \cdot [h_{t-1}, x_t] + b_i)$$

- cell state 에 더해질 수 있는 새로운 후보값 (-1~1)

$$\tilde{C}_t = \tanh(W_C \cdot [h_{t-1}, x_t] + b_C)$$

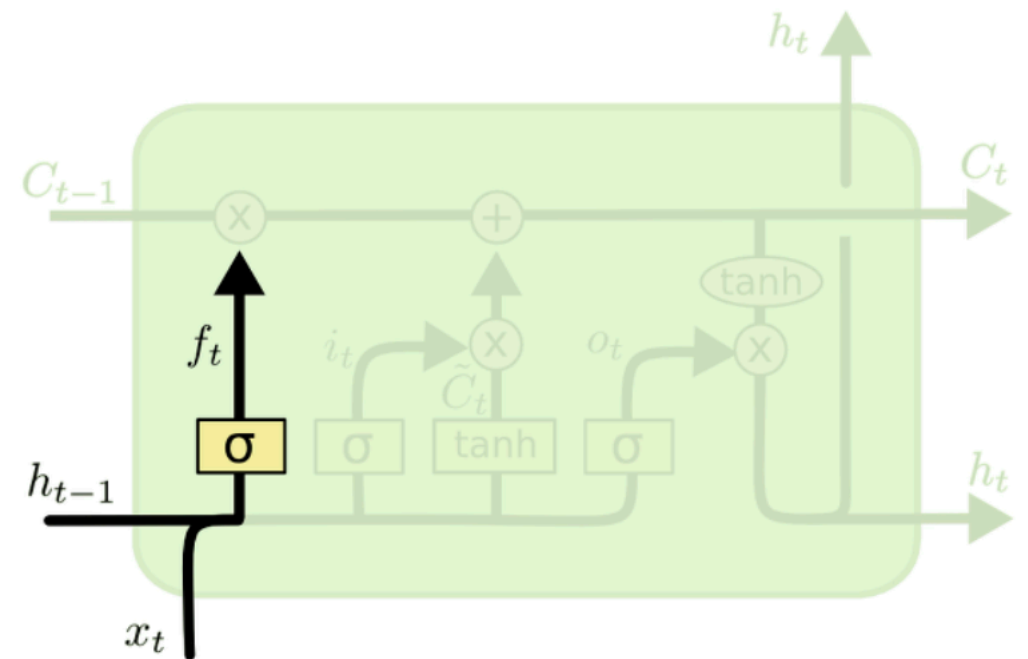
▶ 출력

- Output gate layer: 어떤 출력값을 출력할지 결정 (0~1)

$$o_t = \sigma(W_o [h_{t-1}, x_t] + b_o)$$

- 실제 출력값

$$h_t = o_t * \tanh(C_t)$$



LSTM components

▶ 저장할 정보 (cell states)

$$C_t = f_t * C_{t-1} + i_t * \tilde{C}_t$$

▶ 새로운 정보

- Forget gate layer: cell state에서 어떤 정보를 버릴지 c(0~1)

$$f_t = \sigma(W_f \cdot [h_{t-1}, x_t] + b_f)$$

- Input gate layer: 어떤 값을 업데이트 할지 결정 (0~1)

$$i_t = \sigma(W_i \cdot [h_{t-1}, x_t] + b_i)$$

- cell state 에 더해질 수 있는 새로운 후보값 (-1~1)

$$\tilde{C}_t = \tanh(W_C \cdot [h_{t-1}, x_t] + b_C)$$


▶ 출력

- Output gate layer: 어떤 출력값을 출력할지 결정 (0~1)

$$o_t = \sigma(W_o [h_{t-1}, x_t] + b_o)$$

- 실제 출력값

$$h_t = o_t * \tanh(C_t)$$



LSTM components

▶ 저장할 정보 (cell states)

$$C_t = f_t * C_{t-1} + i_t * \tilde{C}_t$$

▶ 새로운 정보

- Forget gate layer: cell state에서 어떤 정보를 버릴지 c(0~1)

$$f_t = \sigma(W_f \cdot [h_{t-1}, x_t] + b_f)$$

- Input gate layer: 어떤 값을 업데이트 할지 결정 (0~1)

$$i_t = \sigma(W_i \cdot [h_{t-1}, x_t] + b_i)$$

- cell state 에 더해질 수 있는 새로운 후보값 (-1~1)

$$\tilde{C}_t = \tanh(W_C \cdot [h_{t-1}, x_t] + b_C)$$

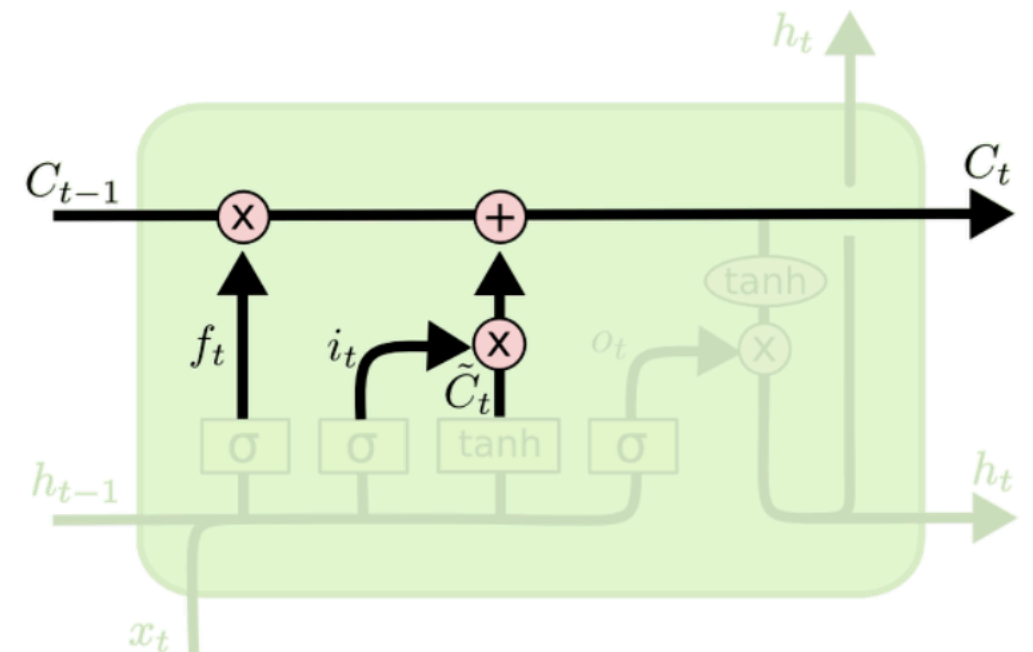
▶ 출력

- Output gate layer: 어떤 출력값을 출력할지 결정 (0~1)

$$o_t = \sigma(W_o [h_{t-1}, x_t] + b_o)$$

- 실제 출력값

$$h_t = o_t * \tanh(C_t)$$



LSTM components

▶ 저장할 정보 (cell states)

$$C_t = f_t * C_{t-1} + i_t * \tilde{C}_t$$

▶ 새로운 정보

- Forget gate layer: cell state에서 어떤 정보를 버릴지 결정 (0~1)

$$f_t = \sigma(W_f \cdot [h_{t-1}, x_t] + b_f)$$

- Input gate layer: 어떤 값을 업데이트 할지 결정 (0~1)

$$i_t = \sigma(W_i \cdot [h_{t-1}, x_t] + b_i)$$

- cell state에 더해질 수 있는 새로운 후보값 (-1~1)

$$\tilde{C}_t = \tanh(W_C \cdot [h_{t-1}, x_t] + b_C)$$


▶ 출력

- Output gate layer: 어떤 출력값을 출력할지 결정 (0~1)

$$o_t = \sigma(W_o [h_{t-1}, x_t] + b_o)$$

- 실제 출력값

$$h_t = o_t * \tanh(C_t)$$



Gated Recurrent Unit (GRU)

► variant of LSTM

- LSTM의 장점을 유지하면서도 계산복잡성을 줄임
- LSTM의 게이트 일부를 생략
 - ✓ input gate와 forget gate를 합침
- cell state와 hidden state (output)을 합침

► Gates of GRU

- update gate ($0 \sim 1$): $z_t = \sigma(W^{(z)}x_t + U^{(z)}h_{t-1})$

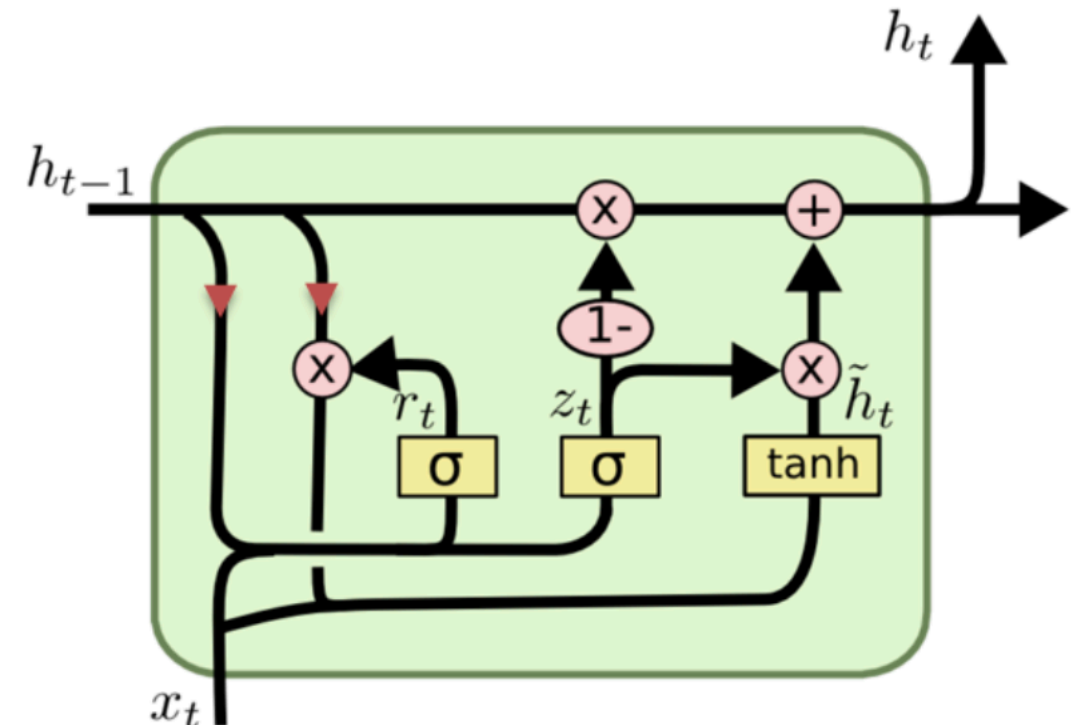
✓ 현재시점 정보와 과거 정보를 어떻게 합칠지

$$\tilde{h}_t = \tanh(Wx_t + r_t \odot Uh_{t-1})$$

- Reset gate ($0 \sim 1$): $r_t = \sigma(W^{(r)}x_t + U^{(r)}h_{t-1})$

✓ 현재시점 정보에 과거 정보 얼마나 반영할지

$$h_t = z_t \odot h_{t-1} + (1 - z_t) \odot \tilde{h}_t$$



Applications of RNNs

- ▶ Machine Translation
- ▶ Image captioning
- ▶ Question Answering

Machine Translation

▶ Input and output data

- 입력데이터: source language에서의 한 문장
- 출력데이터: target language에서의 한 문장

▶ 문제점

- 완전한 단어 matching 이 이루어지기 어려움
- 문장의 길이가 다름

▶ 해결책

- Encoder-decoder scheme
- encoder로 source language의 입력이 다 끝난 이후에 target language의 문장이 나오도록 구성

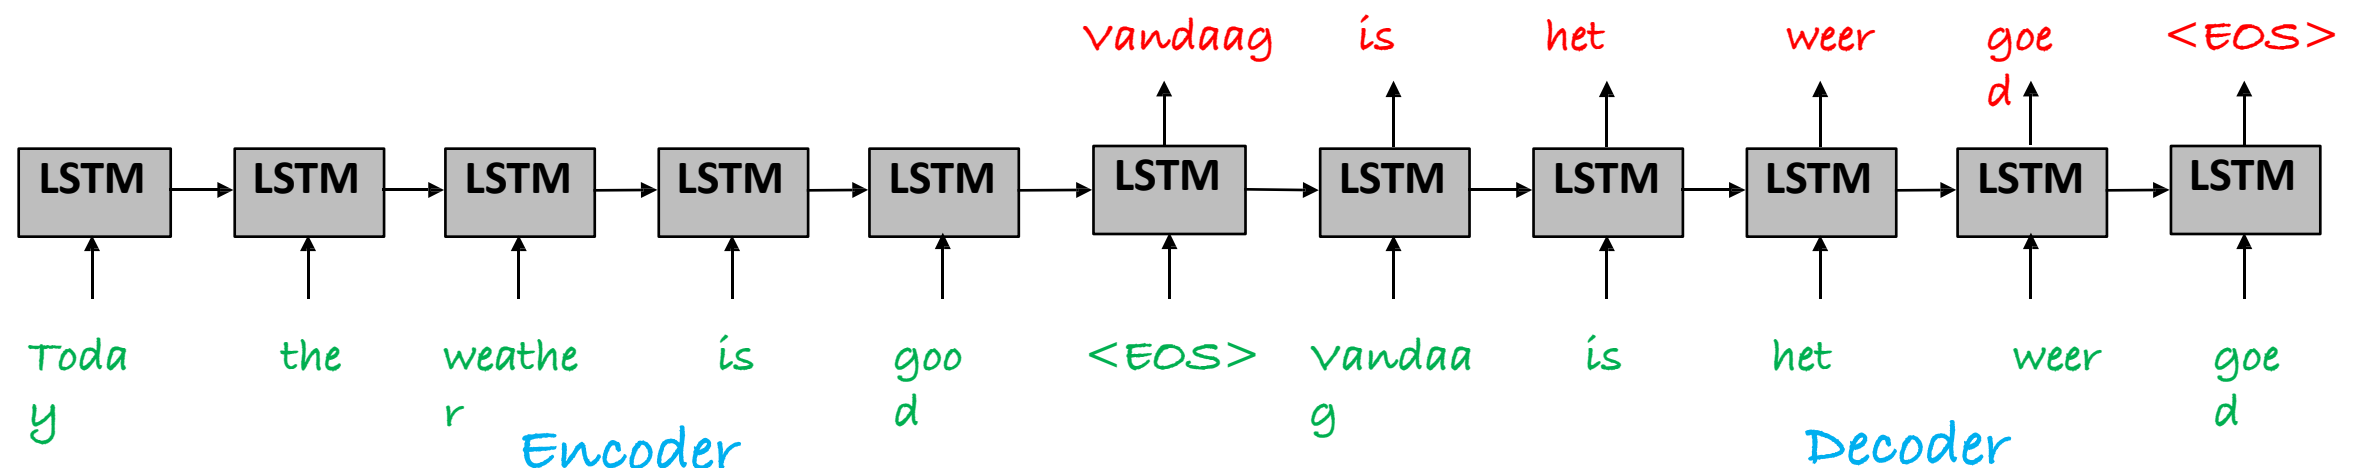


Image captioning

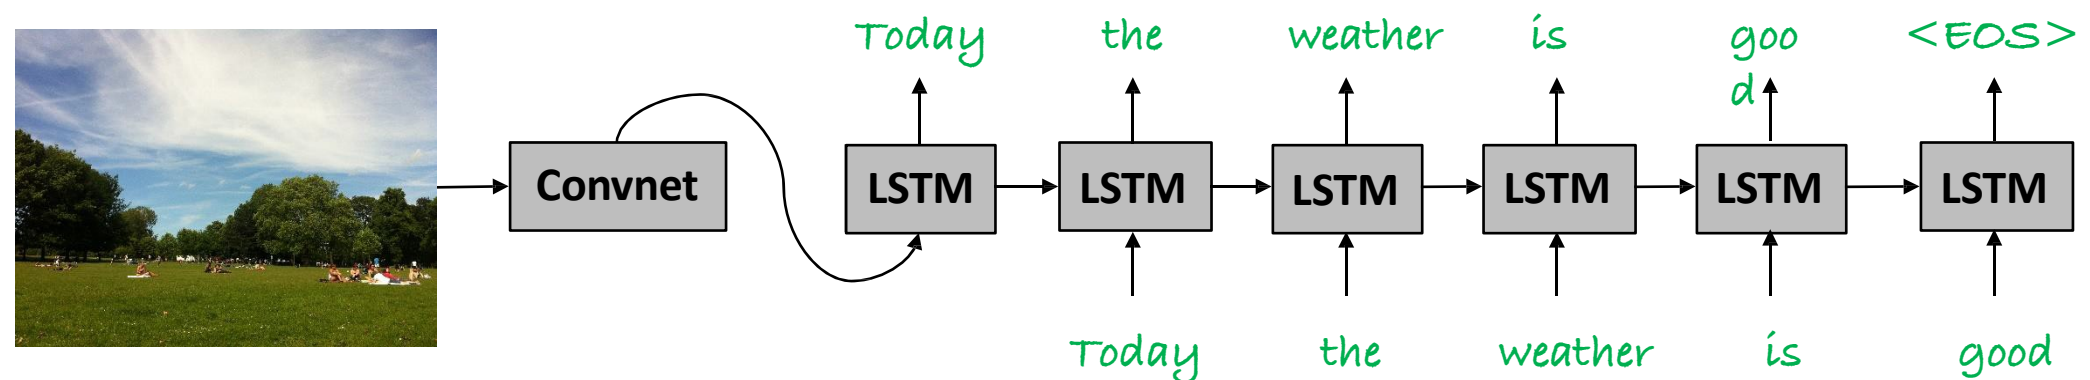
▶ An image is a thousand words, literally!

▶ Input and output data

- 입력데이터: 이미지
- 출력데이터: 문장
- Machine translation과 유사

▶ Network structure

- Encoder 부분을 convnet 으로 대체함
- Decoder part 는 translator 와 동일



Question Answering

► Bleeding-edge research, no real consensus

- very interesting open, research problem
- Machine translation 과 유사
- Encoder-decoder paradigm

► Input data and output data

- 입력데이터: 질문 문장
- 출력데이터: 대답 문장

► Question answering with images

- what has been working so far is to add the image only in the beginning

Q: John entered the living room, where he met Mary. She was drinking some wine and watching a movie. What room did John enter?

A: John entered the living room.



Q: what are the people playing?

A: They play beach football